

[What we do](#) [Blog](#) [Contact us](#) [About](#)[All Posts](#)[Apache Kafka](#)[Google Cloud Platform](#)[BigQuery](#)[Connect](#)[More ▾](#)[Log in / Sign up](#)

Angelo Cesaro Dec 23, 2019 3 min read



Confluent Platform VS Apache Kafka: What are the differences?

Updated: Jan 27

Apache Kafka is an open source software originally created at LinkedIn in 2011. It's used by companies like LinkedIn, Uber, Twitter and more than one-third of all Fortune 500 companies use Apache Kafka. It provides a framework for collecting, reading and analysing streaming data.

Apache Kafka works as a distributed publish-subscribe messaging system. In the Publish-Subscribe system, message producers are called publishers and message consumers are called subscribers. Kafka messages are persisted on the disk and replicated within the cluster to prevent data loss.

Hi there! Do you need help?



Following are a few benefits of Kafka:

- **Reliability:** Kafka is distributed, partitioned, replicated and fault tolerant.
- **Scalability:** Kafka messaging system scales easily without down time by adding additional nodes
- **Durability:** Kafka uses Distributed commit log which means messages persist on disk as fast as possible, hence it is durable..
- **Performance:** Kafka is capable of handling high-volume data. Also, it's able to support message throughput of thousands of messages per second. It is capable of handling these messages with a very low latency.

Now let's see what is Confluent.

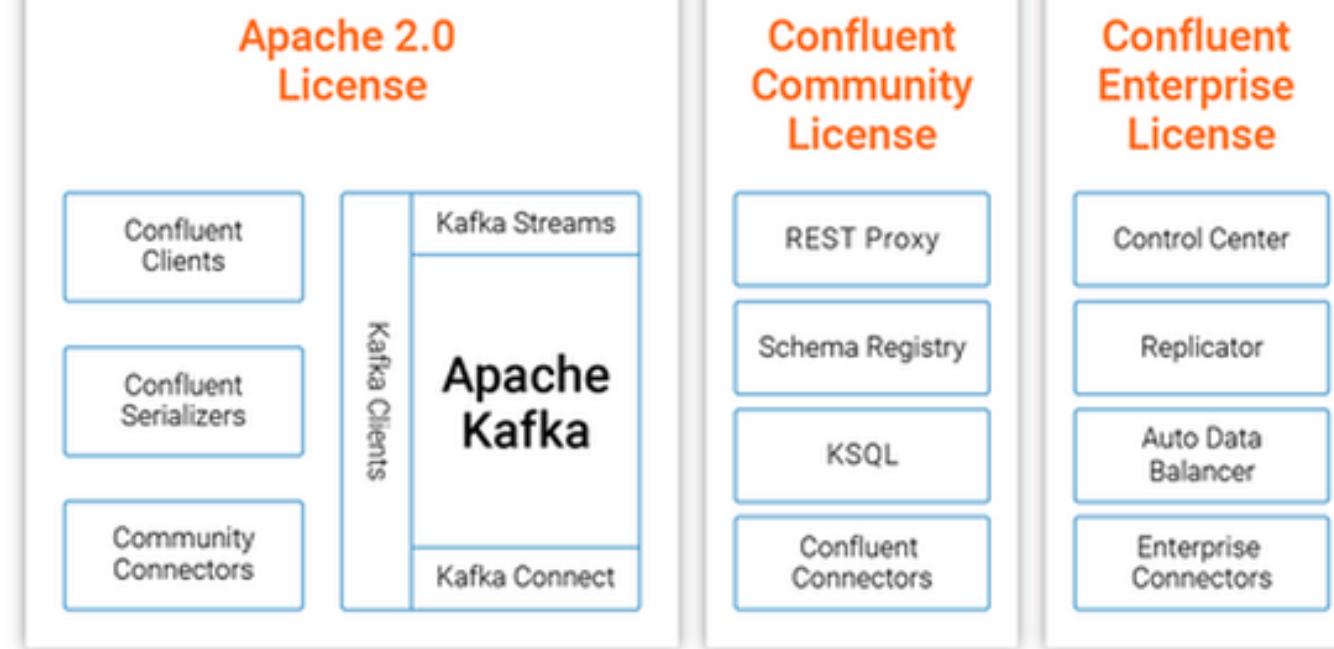
Confluent inc is a company founded in 2004 by the original creators of Apache Kafka. It develops Confluent Platform and provides support, consulting, and training.

Confluent Platform includes **Apache Kafka** and additional (optional) add-ons.

There are open source and community license components of Confluent Platform that are free to use: Rest Proxy, Schema Registry, KSQL and some connectors.

Then there are other components of Confluent Platform that are not free. These components are under the Enterprise license: Control Center, Replicator, Auto Data Balancer and Enterprise Connectors. The Confluent Platform Enterprise components are free forever on single Kafka broker or Free 30 day trial on unlimited Kafka brokers

Hi there! Do you need help?



Confluent Platform Community Components:

- REST Proxy

Provides access to Apache Kafka via HTTP protocol and it operates as a REST API. So you don't need to use the native Kafka protocol to produce messages, consumers, and topics.

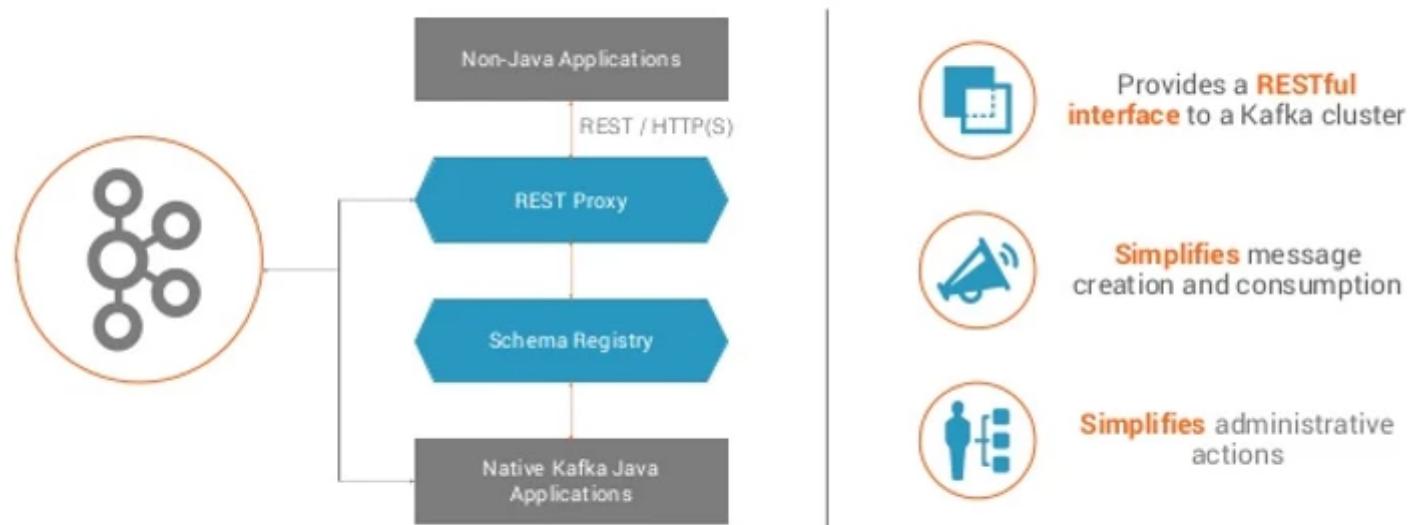
REST API Web API

Hi there! Do you need help?



<https://docs.confluent.io/current/kafka-rest/quickstart.html>

Confluent REST Proxy



confluent

Talking to Non-native Kafka Apps and Outside the Firewall

10

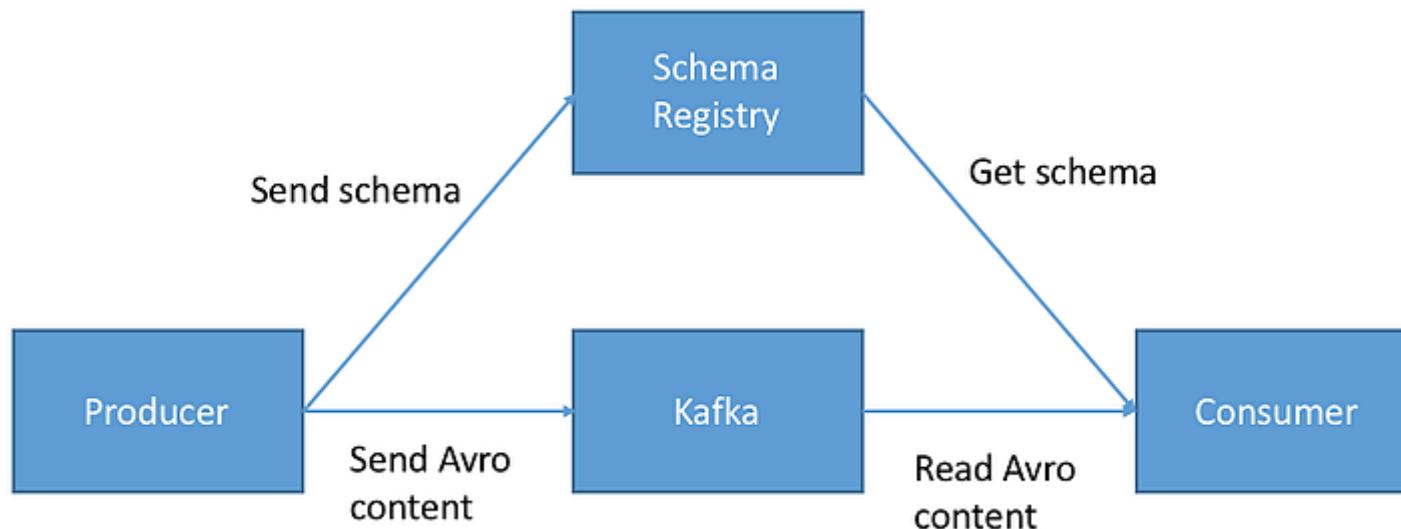
- Schema Registry

Usually Apache Kafka doesn't make data verification. It just takes bytes from input and redistributes them in output. So if the producer sends bad data, or a file data format changes, then the consumers break. The Schema Registry

Hi there! Do you need help?

[What we do](#)[Blog](#)[Contact us](#)[About](#)

schemas are stored in the Schema Registry.



- [KSQL](#)

It's the SQL engine for Apache Kafka to do real-time analysis of data in a topic. It allows you to write SQL queries to explore Kafka topic data easily.

Hi there! Do you need help?



- Pre-built Connectors

Connectors are software that write data from an external data system into Kafka and from Kafka into an external data system. Confluent provides a list of additional supported connectors: Big Query Connector, ElasticSearch Connector, Amazon S3 Connector, Azure Blob Storage Connector, Cassandra Sink Connector, etc. You can see the entire list here: <https://www.confluent.io/hub/> You can use these connectors to import and export data from some of the most commonly used data systems.

Hi there! Do you need help?



What we do Blog Contact us About

Confluent



Kafka Connect
Elasticsearch
Confluent, Inc.

8 68.43%

Confluent



Kafka Connect HDFS
Confluent, Inc.

11 43.26%

Confluent



Kafka Connect JDBC
Confluent, Inc.

4 99.02%

Confluent



Kafka Connect S3
Confluent, Inc.

4 97.77%

Confluent



Confluent Kafka
Replicator
Confluent, Inc.

3 50%

Confluent



Kafka Connect IBM
MQ
Confluent, Inc.

1 66.67%

Confluent



Kafka Connect
ActiveMQ
Confluent, Inc.

0 50%

Confluent



Kafka Connect JMS
Confluent, Inc.

2 100%

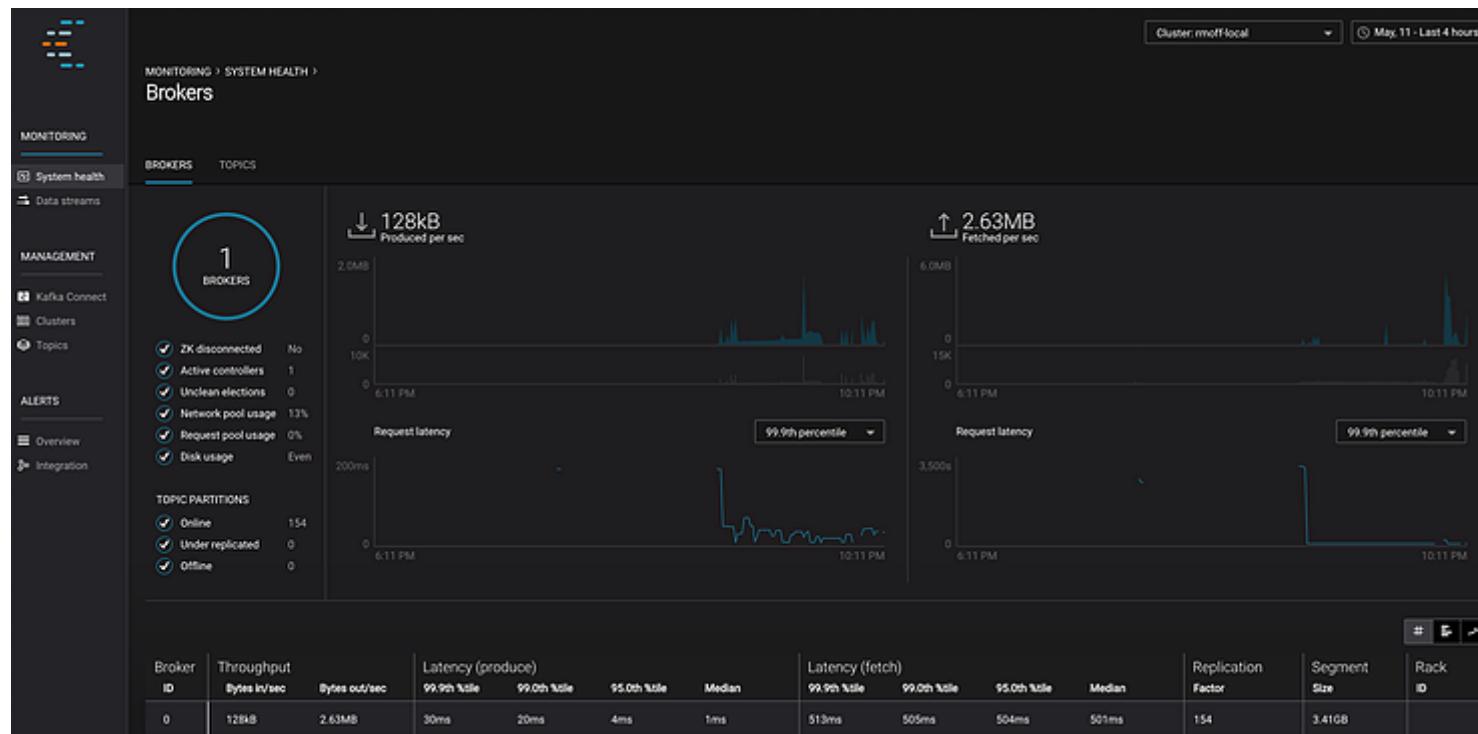
Confluent Platform Enterprise Components:

- [Control Center](#)

Hi there! Do you need help?



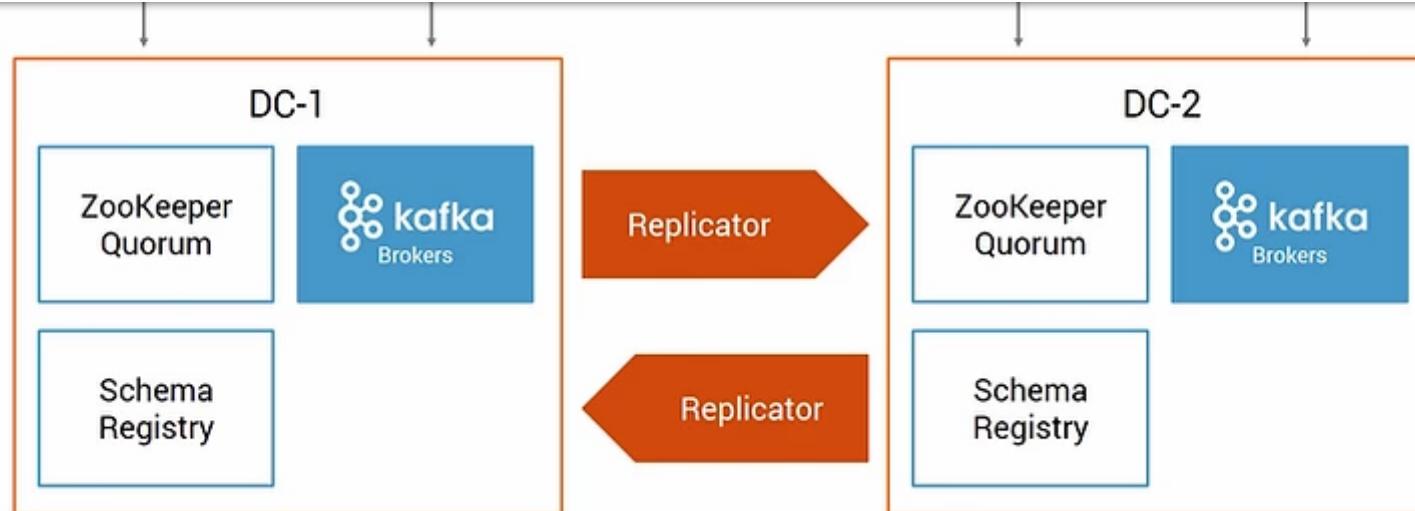
What we do Blog Contact us About



- [Replicator](#)

It allows to replicate data from one cluster to another cluster. So it provides disaster recovery protection to avoid data loss.

Hi there! Do you need help?

[What we do](#)[Blog](#)[Contact us](#)[About](#)

- [Auto Data Balancer](#)

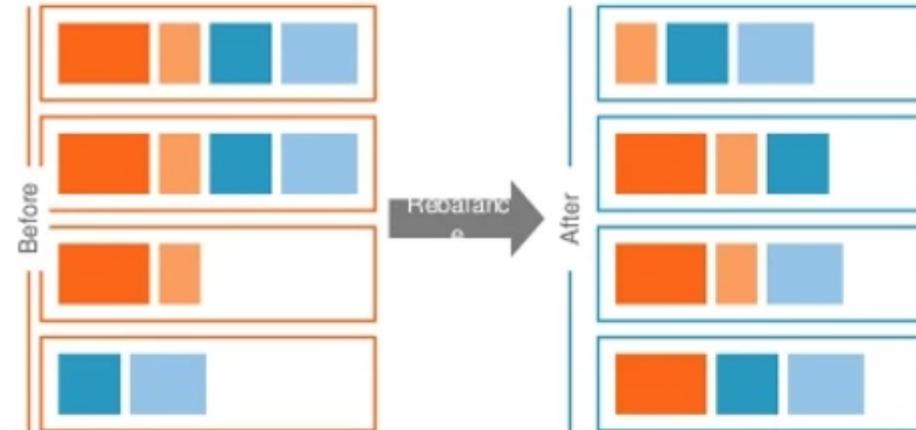
Confluent Auto Data Balancer monitors your cluster for number of nodes, size of partitions, number of partitions and number of leaders within the cluster, and shifts data automatically to create an even workload across your cluster.

Hi there! Do you need help?

[What we do](#)[Blog](#)[Contact us](#)[About](#)

Dynamically move partitions to optimize resource utilization and reliability

- Easily add and remove nodes from your Kafka cluster
- Rack aware algorithm rebalances partitions across a cluster
- Traffic from balancer is throttled when data transfer occurs



-confluent

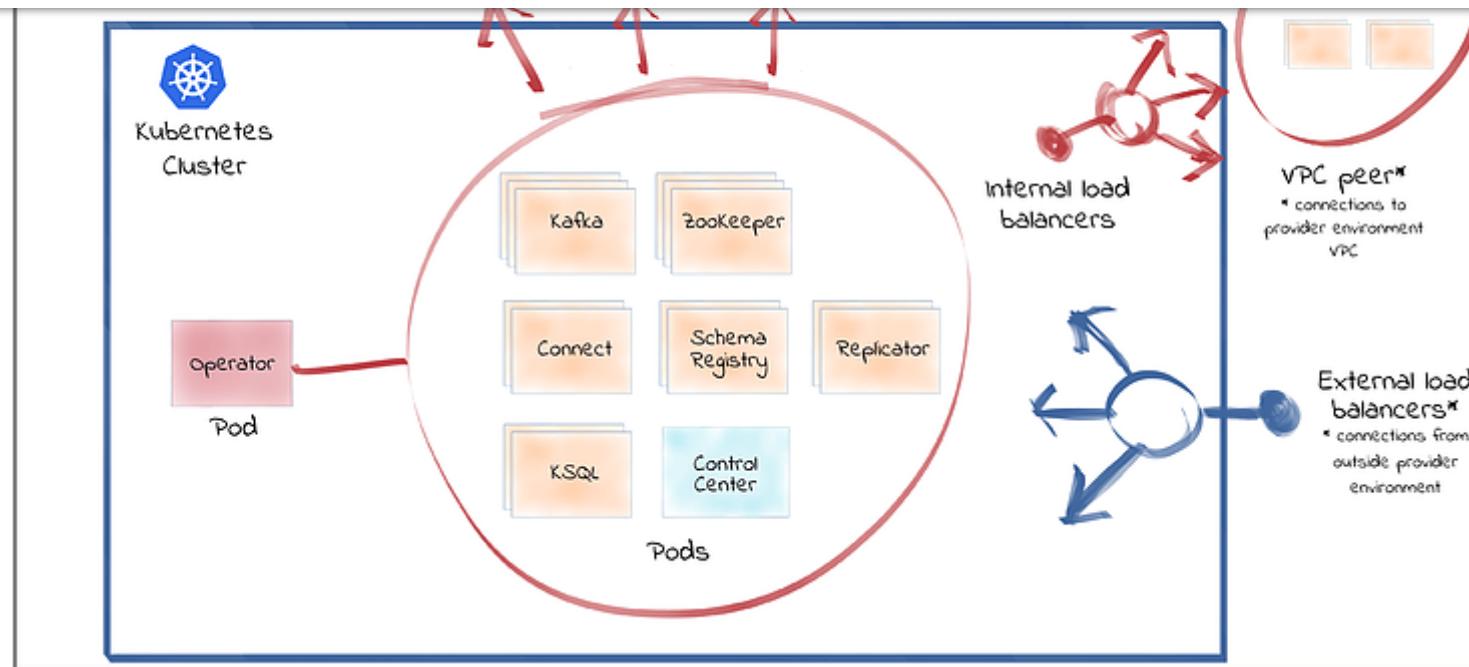
| 44

- Operator

The Confluent Operator provides deployment and management automation for Confluent Platform on Kubernetes (using Helm) including Apache Kafka, Zookeeper, Schema Registry, Connect, Control Center, Replicator and KSQL. Confluent Operator support the following environments: Amazon Elastic Kubernetes Service (EKS), Azure Kubernetes Service (AKS), Google Kubernetes Engine (GKE), OpenShift (3.9 or later), Pivotal Container Service (PKS), Vanilla Kubernetes (1.9 to 1.15) cluster

[Introduction to Confluent Operator for Kubernetes with Viktor Gamov](#)

Hi there! Do you need help?

[What we do](#)[Blog](#)[Contact us](#)[About](#)

Apache Kafka • Connect • Google Cloud Platform

925 views

10

Hi there! Do you need help?



What we do Blog Contact us About

Recent Posts

See All



Apache Kafka and BigQuery connector with batch loads

👁 119

6 ❤

✍ Commenting is disabled on this post.

Hi there! Do you need help?



What we do Blog Contact us About

What we do

Blog

Contact us

About

info@cesaro.io

Cesaro.io LTD

1 Derwent Business Centre

Clarke Street

Derby DE1 2BU, United

Kingdom

Apache, [Apache Kafka](#), [Kafka](#) and associated open source project names are trademark of the [Apache Software Foundation](#)
Company Registration Number: 10823769

[Privacy Policy](#)

© Copyright 2017 - 2020 Cesaro.io LTD - All Rights Reserved

Hi there! Do you need help?